



Fast faceting using roaring bitmaps

Ants Aasma

www.cybertec-postgresql.com



Hello

- ▶ Senior database engineer at Cybertec
- ▶ I focus on: Performance, high availability, core hacking
- ▶ We also do: support, HA setups, DBA-as-a-service, private cloud, general consulting

Problem statement

- ▶ Customer has hundreds of millions of documents.
- ▶ Wants to provide faceted search over these documents.
- ▶ Need accurate counts, or at worst slightly stale counts.
- ▶ Response time below 2 seconds.

What is faceted search

Account Type: Basic Welcome, Esteban Kozak · [Add Connections](#) · [Settings](#) · [Help](#) · [Sign Out](#)

LinkedIn Home Profile Contacts Groups Jobs Inbox (9) More... Search People linkedin

Search 33,778 results for linkedin [Save this search]

Sort by: Relationship View: Basic

Refine By

Location

- ☒ All Locations
- ☐ United States (20762)
- ☐ San Francisco Bay Area (2444)
- ☐ India (2303)
- ☐ United Kingdom (2101)
- ☐ Greater New York City Area (2079)

Show more...

Enter location name

Relationship

- ☒ All LinkedIn Members
- ☐ 1st Connections (192)
- ☐ 2nd Connections (2706)
- ☐ Group Members (860)
- ☐ 3rd + Everyone Else (30390)

Industry

- ☒ All Industries
- ☐ Information Technology and Services (4219)
- ☐ Marketing and Advertising (2709)
- ☐ Staffing and Recruiting (2301)
- ☐ Computer Software (1947)
- ☐ Internet (1782)

Show more...

Enter industry name

Current Company

Past Company

School

Groups

Reid Hoffman
Partner at Greylock
San Francisco Bay Area | Internet
In Common: 120 shared connections 2 shared groups

David Hahn
Director of Product Management, LinkedIn
San Francisco Bay Area | Internet
In Common: 83 shared connections

Candy Chastain Mielke
Consultant at LinkedIn
Colorado Springs, Colorado Area | Internet
In Common: 128 shared connections 2 shared groups

Wade Burgess
Business Leader
Greater Omaha Area | Internet
In Common: 44 shared connections 2 shared groups

Alex Vauthey
Director of Engineering, Monetization at LinkedIn
San Francisco Bay Area | Information Technology and Services
In Common: 103 shared connections

Ruslan Belkin
Sr. Director of Engineering at LinkedIn
San Francisco Bay Area | Computer Software
In Common: 95 shared connections

Elizabeth Reaves
Product Manager at LinkedIn
San Francisco Bay Area | Internet
In Common: 92 shared connections 4 shared groups

Riccardo Ferretti
Principal Engineer at LinkedIn
San Francisco Bay Area | Computer Software
In Common: 121 shared connections

Shirley Xu
Data Scientist at LinkedIn

Search tools

- [Advanced search](#)
- [Saved Searches](#)

Ads by LinkedIn Members

S&OP Summit
 group
January 20/21 2010,
The Palms Hotel Sales
& Operations Planning
Innovat
[events.linkedin.com](#)
From: Richard Bracchi

Channel Not Performing?

On demand solutions
for channel
performance
optimization.
[www.relayware.com](#)
From: Mike Morgan

What's this?

Example schema for testing

Table "test2.documents"				
Column	Type	Collation	Nullable	Default
id	bigint		not null	
created	timestamp with time zone		not null	
finished	timestamp with time zone			
category_id	bigint		not null	
tags	text[]			
type	public.mimetype			
size	bigint			
title	text			

Indexes:

"documents_category_id_idx" btree (category_id)

```
postgres=# SELECT category_id, ROUND(COUNT(*)/1e6, 3) millions FROM documents GROUP BY ROLLUP (category_id) ORDER BY 2 DESC LIMIT 6;
category_id | millions
```

```
-----+-----
      | 100.000
    24 |  60.812
      |  15.214
    12 |   6.764
      |   3.811
    49 |   2.444
```

Facets and example query

- ▶ Facets:
 - ▶ Created month
 - ▶ Finished month
 - ▶ Category
 - ▶ Type
 - ▶ Size class
- ▶ Query: find all other facet counts when filtered by most popular category

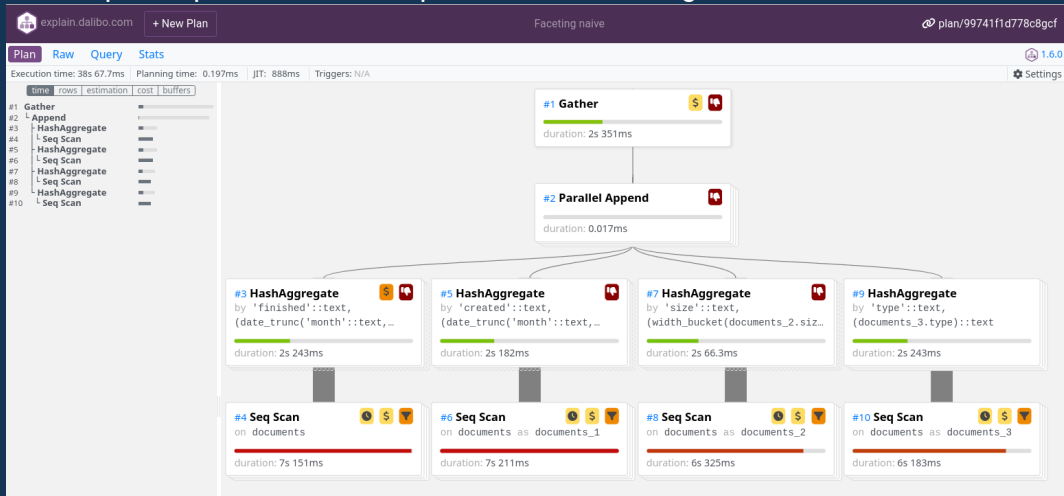


How to implement faceting, simple version

- ▶ Simple option, do a `count(*)` group by for every facet.

```
(SELECT 'created' AS facet_name, date_trunc('month', created)::text AS facet_value,  
      COUNT(*) AS cardinality  
FROM documents WHERE category_id = 24 GROUP BY 1, 2)  
  UNION ALL  
(SELECT 'finished', date_trunc('month', finished)::text, COUNT(*)  
FROM documents  WHERE category_id = 24 GROUP BY 1, 2)  
  UNION ALL  
(SELECT 'type', type::text, COUNT(*)  
FROM documents WHERE category_id = 24 GROUP BY 1, 2)  
  UNION ALL  
(SELECT 'size', width_bucket(size, array[0,1000,5000,10000,50000,100000,500000])::text,  
      COUNT(*)  
FROM documents WHERE category_id = 24 GROUP BY 1, 2);
```

38s: <https://explain.dalibo.com/plan/99741f1d778c8gcf>



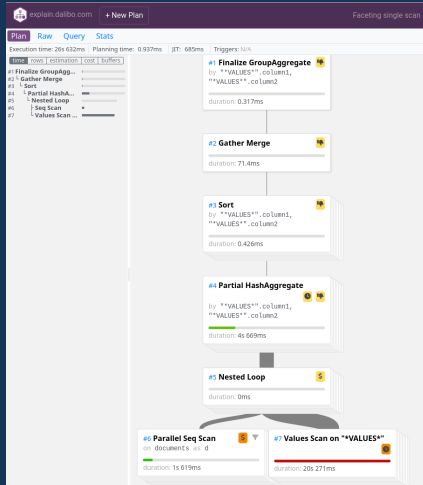
Getting rid of multiple scans

```
SELECT facet_name, facet_value, COUNT(*) cardinality
FROM documents d, LATERAL (VALUES
    ('created', date_trunc('month', created)::text),
    ('finished', date_trunc('month', finished)::text),
    ('type', type::text),
    ('size', width_bucket(size, array[0,1000,5000,10000,50000,100000,500000]
) t(facet_name, facet_value)
WHERE category_id = 24 GROUP BY 1, 2
```



Results

27s: <https://explain.dalibo.com/plan/befh3292ad7cdb12>



More parallelism?

```
ALTER TABLE documents SET ( parallel_workers = 23 );
```

More parallelism?

```
ALTER TABLE documents SET ( parallel_workers = 23 );
```

▶ Time: 18261.594 ms (00:18.262)

Is this even possible?

- ▶ Need to process 60M rows in 2s.
- ▶ For each matching row and each facet:
 - ▶ Calculate the facet value from row
 - ▶ Look up the counter for this facet value and add 1 to it.
- ▶ We have $2s/60e6/4 = 8ns$ of time to do this
- ▶ RAM access takes $\sim 80ns$, L3 cache access is 15ns



What if we invert the problem

- ▶ Random access is really bad for CPUs. Can we turn the problem into a sequential one?
- ▶ What if for each facet and value we store a bitmap of matching documents.
- ▶ Counting matches would then be AND'ing together of two bitmaps and counting the number of bits set.
- ▶ CPU's can do this at $128\text{bits} \times 4\text{GHz} = 512 \text{ Gbit/s/core}$
- ▶ Could calculate number of matches for 10'000 facet values per second per core.



Not a new idea

- ▶ Concept is also known as bitmap indexes.
- ▶ Ideally we would want to use different approaches for storing 60M values and 60 values.
- ▶ A popular and fast implementation is Roaring bitmaps. (roaringbitmap.org)
 - ▶ Used by others too including Solr/Elasticsearch, Pinot, Hive, ClickHouse.
- ▶ Good news, there is a postgres extension wrapping it:
 - ▶ github.com/ChenHuajun/pg_roaringbitmap



What are roaring bitmaps

- ▶ Compressed storage of a set of integers (e.g. document id)
- ▶ Provides fast SIMD optimized operations for intersection (AND), union (OR) and cardinality (COUNT).
- ▶ Divides values into 16bit ranges (65'536 values per range) called containers.
- ▶ First layer is a sorted list of containers that have values, with starting value, type and pointer for each one.
- ▶ Three types of containers:
 - ▶ Array of 16bit values (up to 4'096 values)
 - ▶ 8KB bitmap
 - ▶ Run length encoded bitmap



pg_roaringbitmap extension

- ▶ Introduces a new roaringbitmap datatype for storing bitmaps in the database
- ▶ Some useful functions for us:
 - ▶ `rb_build_agg(int)` -> roaringbitmap
 - ▶ `rb_and(roaringbitmap, roaringbitmap)` -> roaringbitmap
 - ▶ `rb_or(roaringbitmap, roaringbitmap)` -> roaringbitmap
 - ▶ `rb_and_agg(roaringbitmap)` -> roaringbitmap
 - ▶ `rb_cardinality(roaringbitmap)` -> int

The concept

- ▶ Lets build and maintain a second datastructure of roaring bitmaps for each facet value.
- ▶ Conceptually an index, but we are doing it in “userspace”
- ▶ We plan to be modifying data, but mostly only recent rows.
 - ▶ Store bitmaps as chunks so we can update only modified chunks.
 - ▶ Capture modifications into a “delta” table and batch apply changes.
- ▶ Add some functions to make querying easier



Building the index

```
CREATE TABLE documents_facets AS
SELECT facet_name, facet_value, id >> 20 chunk_id, rb_build_agg((id & ((1<<20)-1))::int4)
FROM documents d, LATERAL (VALUES
    ('created', date_trunc('month', created)::text),
    ('finished', date_trunc('month', finished)::text),
    ('category_id', category_id::text),
    ('type', type::text),
    ('size', width_bucket(size, array[0,1000,5000,10000,50000,100000,500000]))::text)
) t(facet_name, facet_value) GROUP BY 1, 2, 3;
```

Time: 37334.313 ms (00:37.334)



Building the index

```
CREATE TABLE documents_facets AS
SELECT facet_name, facet_value, id >> 20 chunk_id, rb_build_agg((id & ((1<<20)-1))::int4)
FROM documents d, LATERAL (VALUES
    ('created', date_trunc('month', created)::text),
    ('finished', date_trunc('month', finished)::text),
    ('category_id', category_id::text),
    ('type', type::text),
    ('size', width_bucket(size, array[0,1000,5000,10000,50000,100000,500000]))::text)
) t(facet_name, facet_value) GROUP BY 1, 2, 3;
```

Time: 37334.313 ms (00:37.334)

table	table_size	indexes_size
documents	20 GB	661 MB
documents_facets	216 MB	848 kB

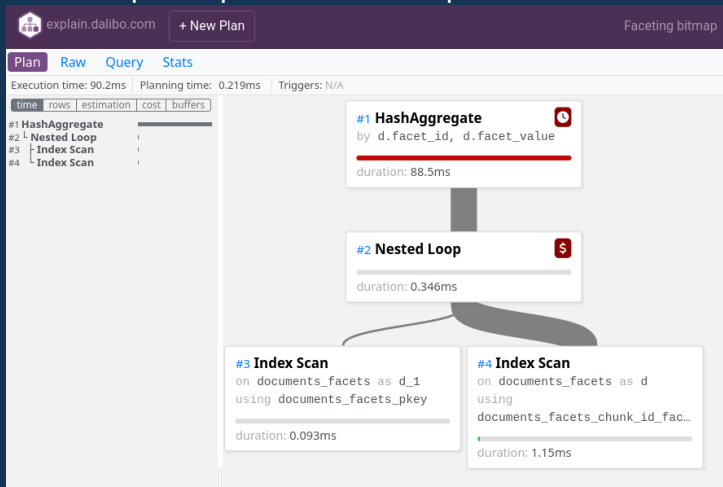
Querying

```
WITH lookup AS (  
    SELECT chunk_id, postinglist FROM documents_facets d  
    WHERE facet_name = 'category_id' AND facet_value = '24'  
)  
SELECT facet_id, facet_value,  
       sum(rb_and_cardinality(lookup.postinglist, d.postinglist)) count  
FROM lookup JOIN documents_facets d USING (chunk_id)  
WHERE facet_id IN (1,2,4,5)  
GROUP BY 1,2;
```



Results

90ms: <https://explain.dalibo.com/plan/71914e60dc7ab324>



Updating the index

- ▶ Don't want to rewrite multiple big bitmaps for each inserted row.
- ▶ Create a delta table with structure (facet_id, facet_value, id, delta)
- ▶ A trigger on main table that inserts into delta table +1 or -1 for each changed facet.
- ▶ Periodically aggregate deltas and merge into main index.



The extension

- ▶ Started an extension called pgfaceting
 - ▶ github.com/cybertec-postgresql/pgfaceting
- ▶ SQL only, but requires pg_roaringbitmap
- ▶ Open source, help appreciated
- ▶ API is still WIP

What is there

- ▶ Automatic creation and population of facet and delta tables.
- ▶ Facet types:
 - ▶ plain
 - ▶ date_trunc
 - ▶ bucket_facet
- ▶ Delta trigger generation.
- ▶ Query generation facility

Example usage

```
SELECT faceting.add_faceting_to_table(  
  'documents',  
  key => 'id',  
  facets => array[  
    faceting.datetrunc_facet('created', 'month'),  
    faceting.datetrunc_facet('finished', 'month'),  
    faceting.plain_facet('category_id'),  
    faceting.plain_facet('type'),  
    faceting.bucket_facet('size',  
      buckets => array[0,1000,5000,10000,50000,100000,500000])  
  ]  
);  
-- Add a cron job that does:  
CALL faceting.run_maintenance();
```

Example usage

```
SELECT * FROM faceting.count_results('documents'::regclass,  
    filters => array[row('category_id', '24'),  
                     row('type', 'image/jpeg')  
    ]::faceting.facet_filter[]);
```

Future ideas

- ▶ Support id's larger than int4
- ▶ New facet types
 - ▶ array (e.g. `unnest(tags)`)
 - ▶ Tree structured facets
 - ▶ Hierarchical date and time
- ▶ Automatic range queries
- ▶ Use delta tables in queries
- ▶ Full text search



When to not use it

- ▶ Doesn't work well with sparse id spaces like uuid, snowflake.
Also rules out using ctid
- ▶ Tons of unique values (tens to hundreds of thousands)
- ▶ Heavily updated facets

Thanks

► Q & A