

FIXING LOCK CONTENTION IN OLTP

ANTS AASMA

PGUG Estonia meetup 2025

Hello



About me

- Cybertec Lead Database Consultant
- PostgreSQL contributor
- Performance, HA, extensions, etc.



The problem statement



Help, the app is too slow

```
1[|] 2.4% 7[|] 2.5% 13[|] 1.8% 19[|] 1.2%
2[|] 1.8% 8[|] 6.2% 14[|] 1.9% 20[|] 1.9%
3[|] 1.8% 9[|] 2.4% 15[|] 1.9% 21[|] 2.5%
4[|] 1.2% 10[|] 1.2% 16[|] 3.0% 22[|] 2.5%
5[|] 1.8% 11[|] 8.0% 17[|] 2.4% 23[|] 3.7%
6[|] 1.8% 12[|] 1.9% 18[|] 1.9% 24[|] 3.6%
Mem[||||||||| 28.8G/121G Tasks: 606, 4961 thr, 384 kthr; 1 running
Swp[          0K/0K Load average: 1.14 1.60 1.46
Uptime: 5 days, 02:02:43
```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	DISK WRITE	DISK READ	TIME+	Command
2037124	ants	20	0	8533M	67348	63956	S	1.2	0.1	38.88 K/s	0.00 B/s	0:00.28	postgres: ants postgres [local] UPDATE waitin
2037127	ants	20	0	8533M	65396	62008	S	1.2	0.1	72.90 K/s	0.00 B/s	0:00.28	postgres: ants postgres [local] UPDATE waitin
2037134	ants	20	0	8533M	68584	65192	S	1.2	0.1	92.35 K/s	0.00 B/s	0:00.32	postgres: ants postgres [local] UPDATE waitin
2037139	ants	20	0	8533M	68032	64644	S	1.2	0.1	102.07 K/s	0.00 B/s	0:00.31	postgres: ants postgres [local] UPDATE waitin
2037160	ants	20	0	8533M	68872	65488	S	1.2	0.1	111.79 K/s	0.00 B/s	0:00.33	postgres: ants postgres [local] UPDATE waitin
2037164	ants	20	0	8533M	67144	63756	S	1.2	0.1	116.65 K/s	0.00 B/s	0:00.31	postgres: ants postgres [local] UPDATE waitin
2037178	ants	20	0	8533M	67132	63744	S	1.2	0.1	58.32 K/s	0.00 B/s	0:00.30	postgres: ants postgres [local] UPDATE waitin
2037113	ants	20	0	8533M	67000	63596	S	0.6	0.1	24.30 K/s	0.00 B/s	0:00.29	postgres: ants postgres [local] UPDATE waitin
2037115	ants	20	0	8533M	68392	65000	S	0.6	0.1	43.74 K/s	0.00 B/s	0:00.33	postgres: ants postgres [local] UPDATE waitin
2037118	ants	20	0	8533M	66196	62812	S	0.6	0.1	58.32 K/s	0.00 B/s	0:00.28	postgres: ants postgres [local] UPDATE waitin
2037119	ants	20	0	8533M	66536	63148	S	0.6	0.1	43.74 K/s	0.00 B/s	0:00.28	postgres: ants postgres [local] UPDATE waitin
2037120	ants	20	0	8533M	68896	65508	S	0.6	0.1	92.35 K/s	0.00 B/s	0:00.33	postgres: ants postgres [local] UPDATE waitin



Wait events to the rescue

```
postgres=# SELECT wait_event_type, wait_event, count(*)
postgres-# FROM pg_stat_activity
postgres-# WHERE application_name = 'pgbench'
postgres-# GROUP BY 1, 2 ORDER BY 3 DESC;
      wait_event_type |    wait_event    | count
-----+-----+-----+
      Lock          | transactionid |     86
      Lock          | tuple         |     13
      IO            | WalSync       |      1
(3 rows)
```



What's happening

Request A:

```
BEGIN;  
UPDATE acct SET b = b + 15 WHERE id = 1;
```

Request B:

```
-----  
-----
```



What's happening

Request A:

```
BEGIN;  
UPDATE acct SET b = b + 15 WHERE id = 1;  
  
-- waiting for client
```

Request B:

```
BEGIN;  
UPDATE acct SET b = b - 5 WHERE id = 1  
          AND b > 5;  
  
-- Blocks . . .
```



What's happening

Request A:

```
BEGIN;  
UPDATE acct SET b = b + 15 WHERE id = 1;  
  
-- waiting for client  
  
COMMIT;
```

Request B:

```
BEGIN;  
UPDATE acct SET b = b - 5 WHERE id = 1  
    AND b > 5;  
  
-- Blocks . . .  
  
-- Update completes  
  
COMMIT;
```

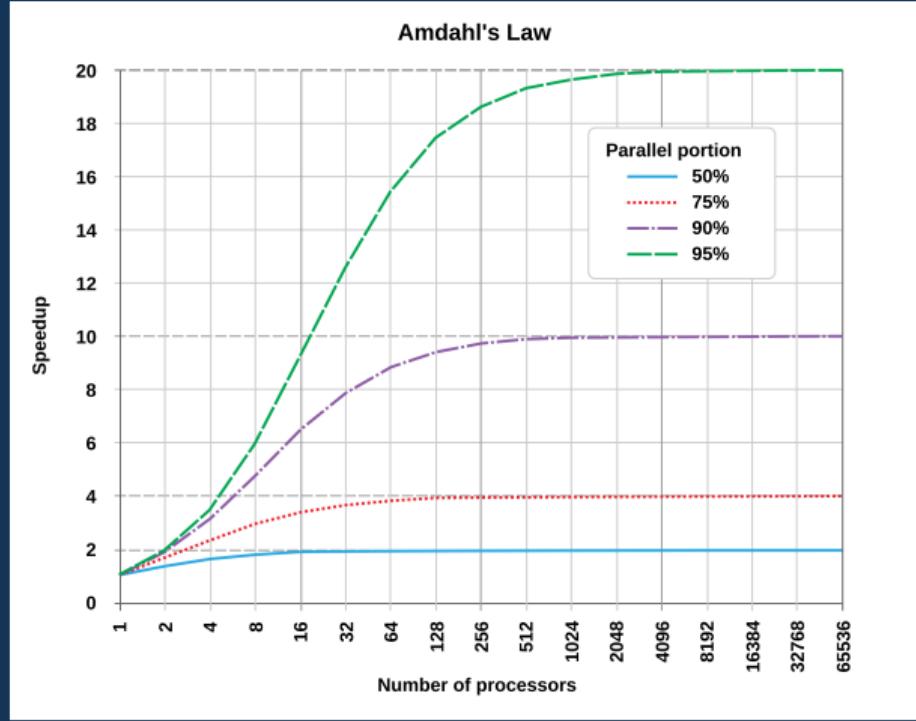


Why this is a big problem

- Double entry accounting requires two accounts for every transaction.
- Some accounts are very hot.
 - “The Bank”
 - “The Payment Processor”
- Throughput limited by update rate on the hot account.



Amdahl's law strikes again



By Daniels220 at English Wikipedia, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=6678551>



The trouble with SQL



Interactive protocol

```
BEGIN;  
UPDATE accounts ... WHERE act = :a1 ...;  
UPDATE accounts ... WHERE act = :a2 ...;  
INSERT INTO transactions ...;  
SELECT ...;  
COMMIT;
```



Interactive protocol

```
BEGIN;  
UPDATE accounts ... WHERE act = :a1 ...;  
UPDATE accounts ... WHERE act = :a2 ...;  
INSERT INTO transactions ...;  
SELECT ...;  
COMMIT;
```

- Database has no idea what will happen next.



Interactive protocol

```
BEGIN;  
UPDATE accounts ... WHERE act = :a1 ...;  
UPDATE accounts ... WHERE act = :a2 ...;  
INSERT INTO transactions ...;  
SELECT ...;  
COMMIT;
```

- Database has no idea what will happen next.
- Every command is a network roundtrip.



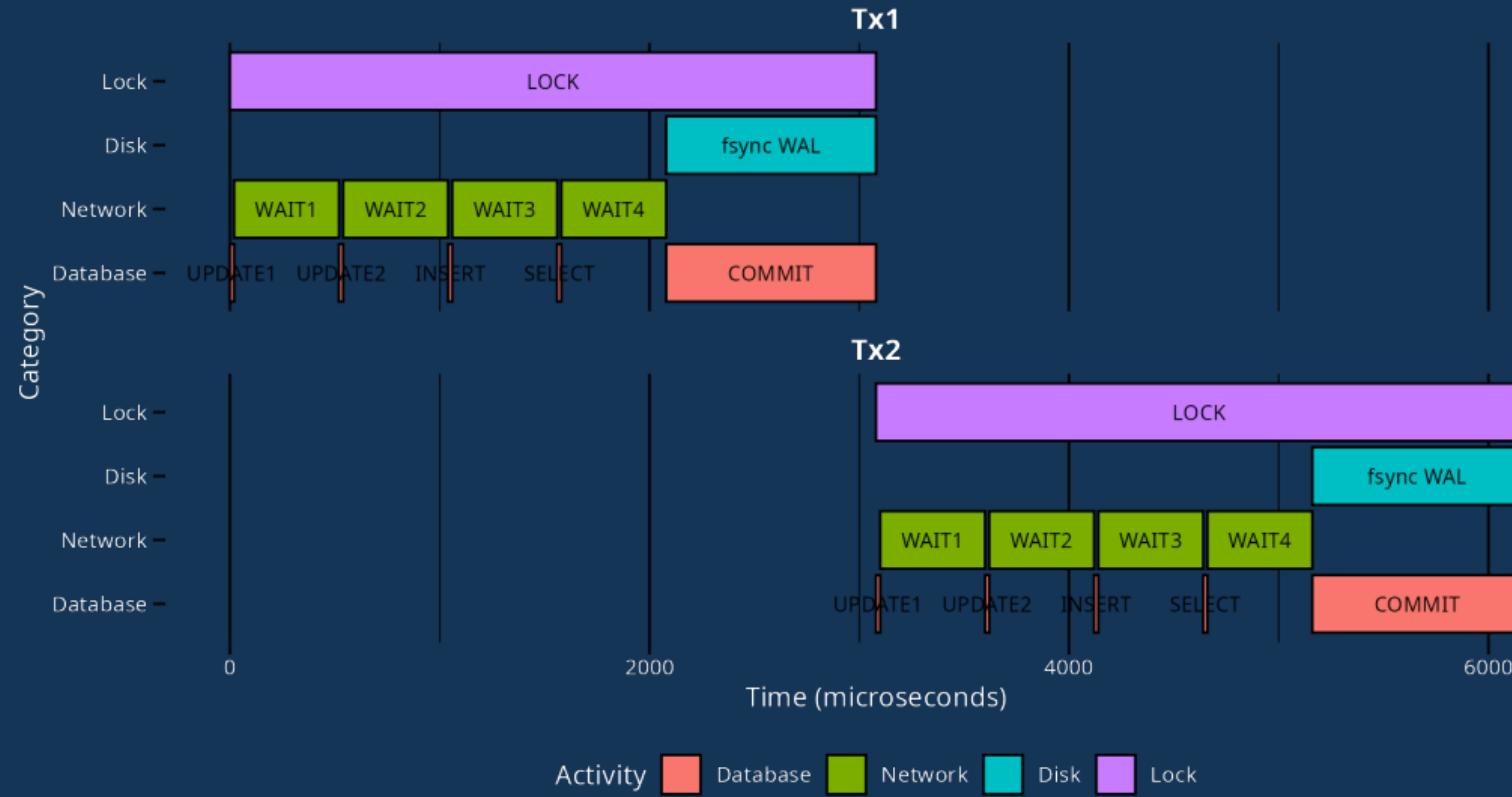
Ballpark numbers

- UPDATE:
 - HOT update using PK: 0.02 ms
- Network roundtrip:
 - Fast network: 0.1 .. 0.2 ms
 - Cloud network: 0.5ms
 - Cloud cross AZ: 0.5 .. 2 ms
 - Cross region: 10 .. 50ms
- Disk roundtrip:
 - Enterprise NVMe: 0.02 ms
 - Good network disk: 0.5ms



How it looks like

Transaction timeline conversational execution



Simple optimizations

- Minimize network roundtrip time.
- Do locking actions late to minimize lock hold time.



Taking the logic to the data



Option 1: CTEs

```
WITH update1 AS (UPDATE ... RETURNING ...),
      update2 AS (UPDATE ... RETURNING ...),
      insert1 AS (INSERT ...
                  SELECT ... FROM update1, update2
                  RETURNING ...),
SELECT ... FROM insert1;
```

Will get quite hairy if the logic is more complicated.



Option 2: Pipeline mode

- Send multiple commands to PostgreSQL in a single roundtrip.
- If any command errors the whole pipeline is cancelled.
- Command must raise errors on failure.
- Very few client drivers support it. (libpq, psycopg)



Option 3: Stored procedures

Write the data logic as a single stored function.



Option 3: Stored procedures

Write the data logic as a single stored function.

```
CREATE FUNCTION transfer(p_src int, p_dest int, p_amount numeric) RETURNS int8 AS $$  
DECLARE new_balance numeric; tx_id int8;  
BEGIN  
    UPDATE accounts SET balance = balance - p_amount  
        WHERE id = p_src RETURNING balance INTO new_balance;  
    IF new_balance < 0 THEN RAISE EXCEPTION 'Not enough funds in %', p_src; END IF;  
  
    UPDATE accounts SET balance = balance + p_amount  
        WHERE id = p_dest RETURNING balance INTO new_balance;  
    IF new_balance < 0 THEN RAISE EXCEPTION 'Not enough funds in %', p_dest; END IF;  
  
    INSERT INTO transactions (src, dest, amount)  
        VALUES (p_src, p_dest, p_amount)  
        RETURNING id INTO tx_id;  
  
    RETURN tx_id;  
END;  
$$ LANGUAGE plpgsql;
```



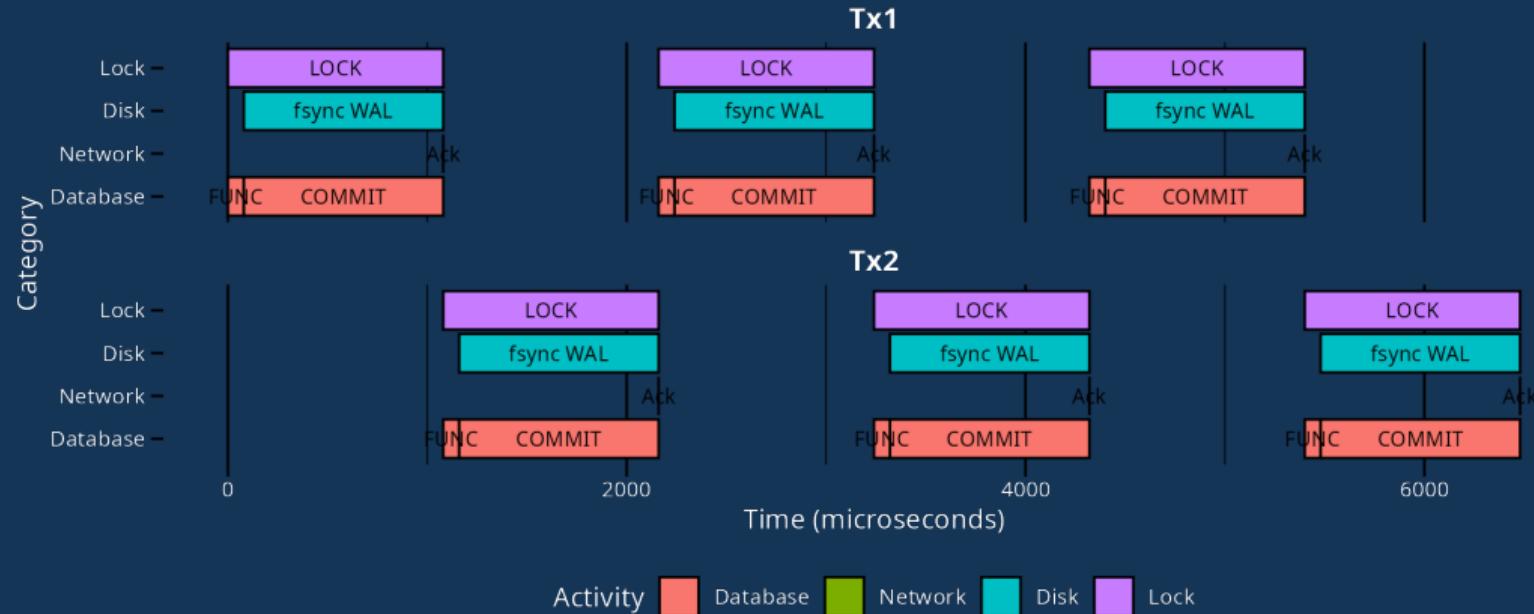
Avoiding all roundtrips

- Send the function call outside explicit transaction.
- Starts implicit transaction in database.
 - If command fails gets rolled back.
 - If command succeeds gets committed automatically.
- Some drivers need special handling for this.
 - psycopg - autocommit=True
 - JDBC - conn.setAutoCommit(true); (default)



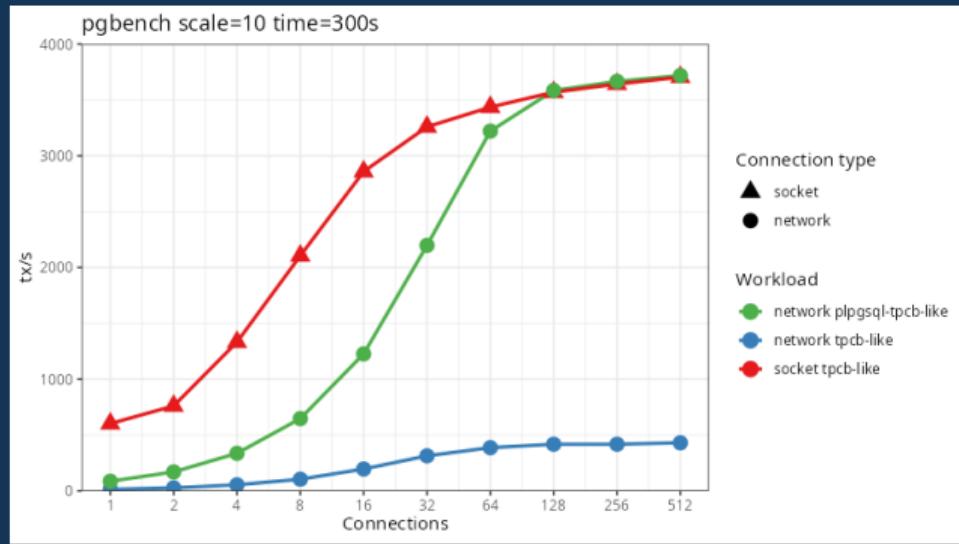
How server side execution looks like

Transaction timeline with server side execution



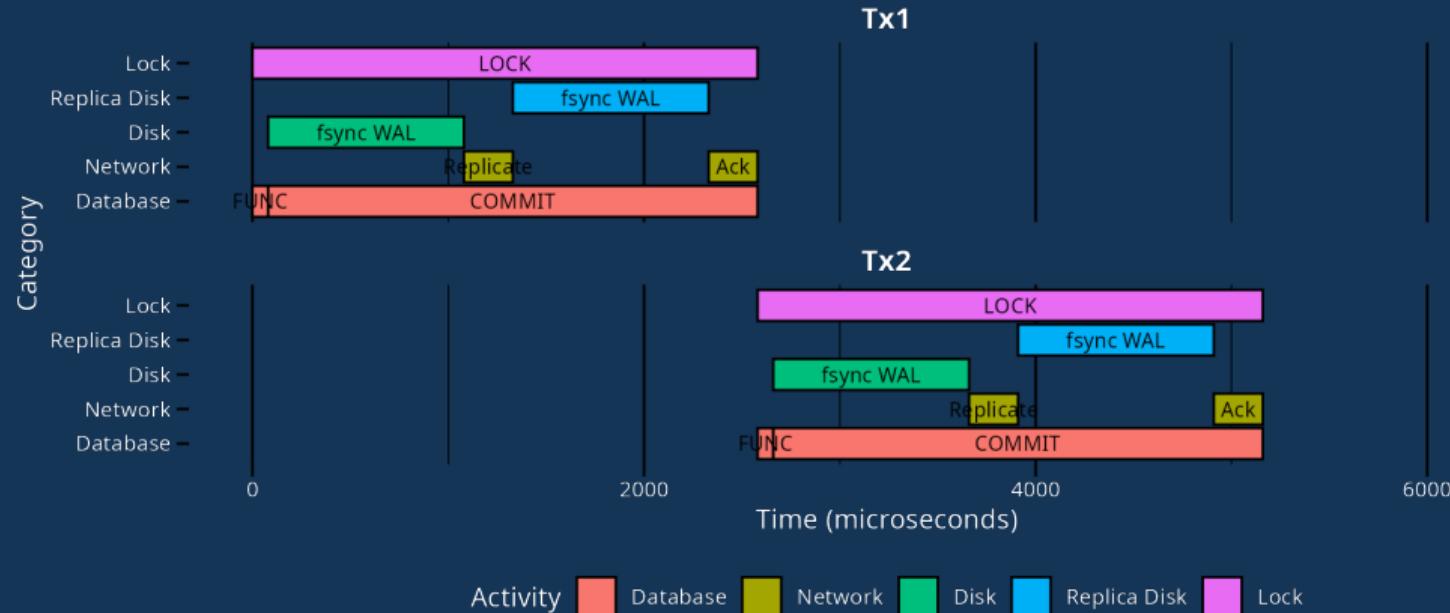
The benefit

Network has 10ms roundtrip latency.



Synchronous replication

Transaction timeline for synchronous replication



Avoiding the commit latency



The naughty option

- `synchronous_commit=off` disables WAL flush.
- From 800 tps to 80'000 tps.
- On any crash/failover you will lose transactions.
- Typically less than `wal_writer_delay = 200ms`
- May be ok if you can detect crashes and replay transactions.



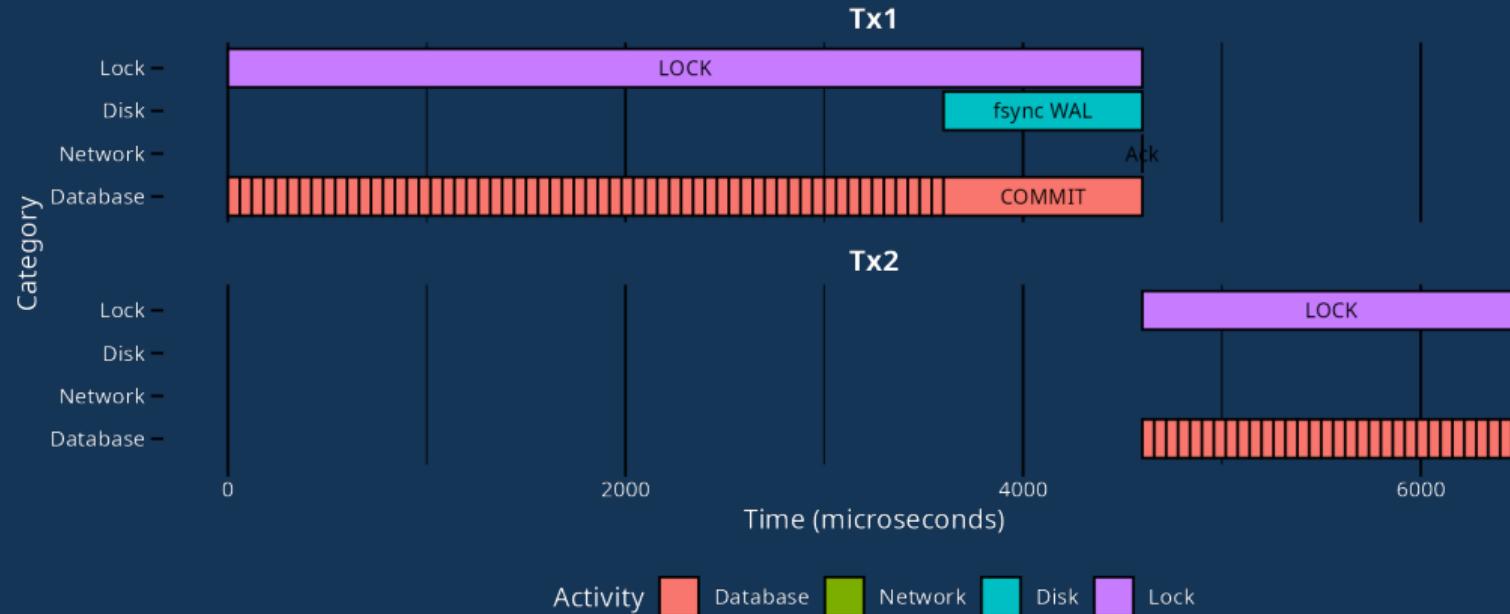
The clever option

- Amortize the durability cost across multiple transactions.
- Collect batch of transactions in application.
- Send to server function for execution.
- Server returns list of result statuses.

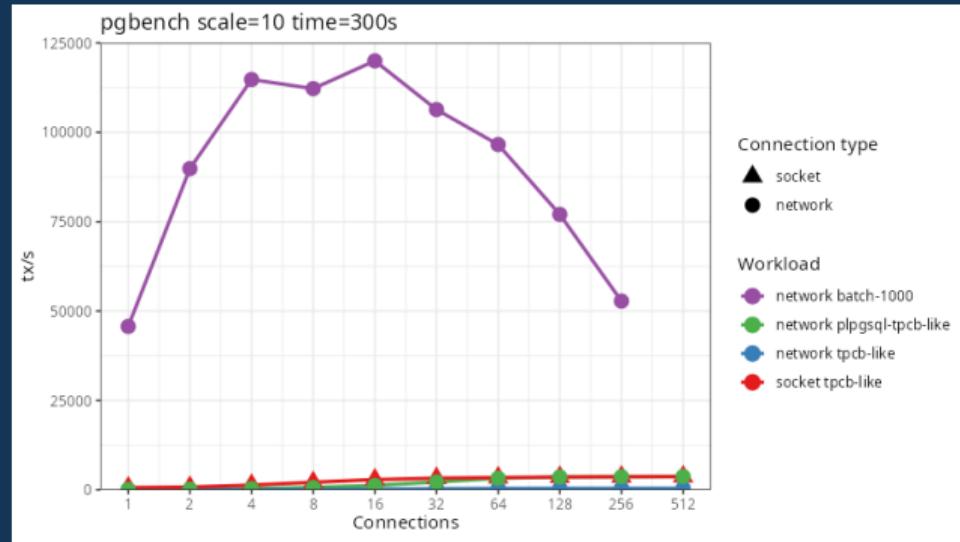


How batching looks like

Transaction Timeline (Faceted by Transaction)



Effect on performance



Downsides

- Complicated work needed to manage batches.
- Need to be careful to limit latency hit.



Transaction isolation



Isolation levels

When two transactions update same row:

READ COMMITTED

1. Block until locker finishes.
2. Look at the new version.
3. Recheck if WHERE condition matches.
4. Update new version.

REPEATABLE READ / SERIALIZABLE

1. Block until locker finishes.
2. If locker committed then raise serialization error.



Contention hell

- Any contention will cause concurrent transaction to fail/retry.
- Extremely inefficient use of database resources.
- Optimistic concurrency control and contention do not mix.
- Can't use even table level locks - requesting a lock acquires a snapshot.



Be careful with correctness

- Read committed allows a single query to see different database states.
- Quite complicated to know if this is correct.



Possible workaround

Session level advisory locks allow snapshot to be released after locking.

```
-- Block until all accounts available  
SELECT pg_advisory_lock(x) FROM unnest(:batch) x ORDER BY x;
```

```
BEGIN ISOLATION LEVEL REPEATABLE READ;
```

```
    -- Do work ...
```

```
COMMIT;
```

```
-- Unlock
```

```
SELECT pg_advisory_unlock_all();
```



Getting rid of contention



Delta tables

- Calculate account balance based on transaction history.
- Only usable when transactions are not dependent.



Double booking problem

Two tables.

- flight.seats is total available seats.
- booking.booked is number of seats booked.

Starting condition: only 2 free seats on a flight.

Try to book 2 seats for A:

```
BEGIN;
  SELECT f.seats - SUM(b.booked) available
    FROM flight f JOIN booking b USING (flight_id)
   WHERE flight_id = 123 GROUP BY flight_id;

  IF available > 2 THEN
    INSERT INTO booking VALUES (123, 'tx A', 2);
  END IF;
COMMIT;
```

Try to book 2 seats for B:

```
BEGIN;
  SELECT f.seats - SUM(b.booked) available
    FROM flight f JOIN booking b USING (flight_id)
   WHERE flight_id = 123 GROUP BY flight_id;

  IF available > 2 THEN
    INSERT INTO booking VALUES (123, 'tx B', 2);
  END IF;
COMMIT;
```



Double booking problem

Repeatable read allows for double booking.
Serializable will fail with serialization error.



Sharded accounts

- Usually hot accounts are not close to empty.



Sharded accounts

- Usually hot accounts are not close to empty.
- Split hot accounts into multiple pieces.



Sharded accounts

- Usually hot accounts are not close to empty.
- Split hot accounts into multiple pieces.
- Balance updates between sub-accounts.



Sharded accounts

- Usually hot accounts are not close to empty.
- Split hot accounts into multiple pieces.
- Balance updates between sub-accounts.
- Handle case when one sub-account is empty.



Sharded accounts

- Usually hot accounts are not close to empty.
- Split hot accounts into multiple pieces.
- Balance updates between sub-accounts.
- Handle case when one sub-account is empty.
- Total balance can be calculated on demand.



Science fiction future features



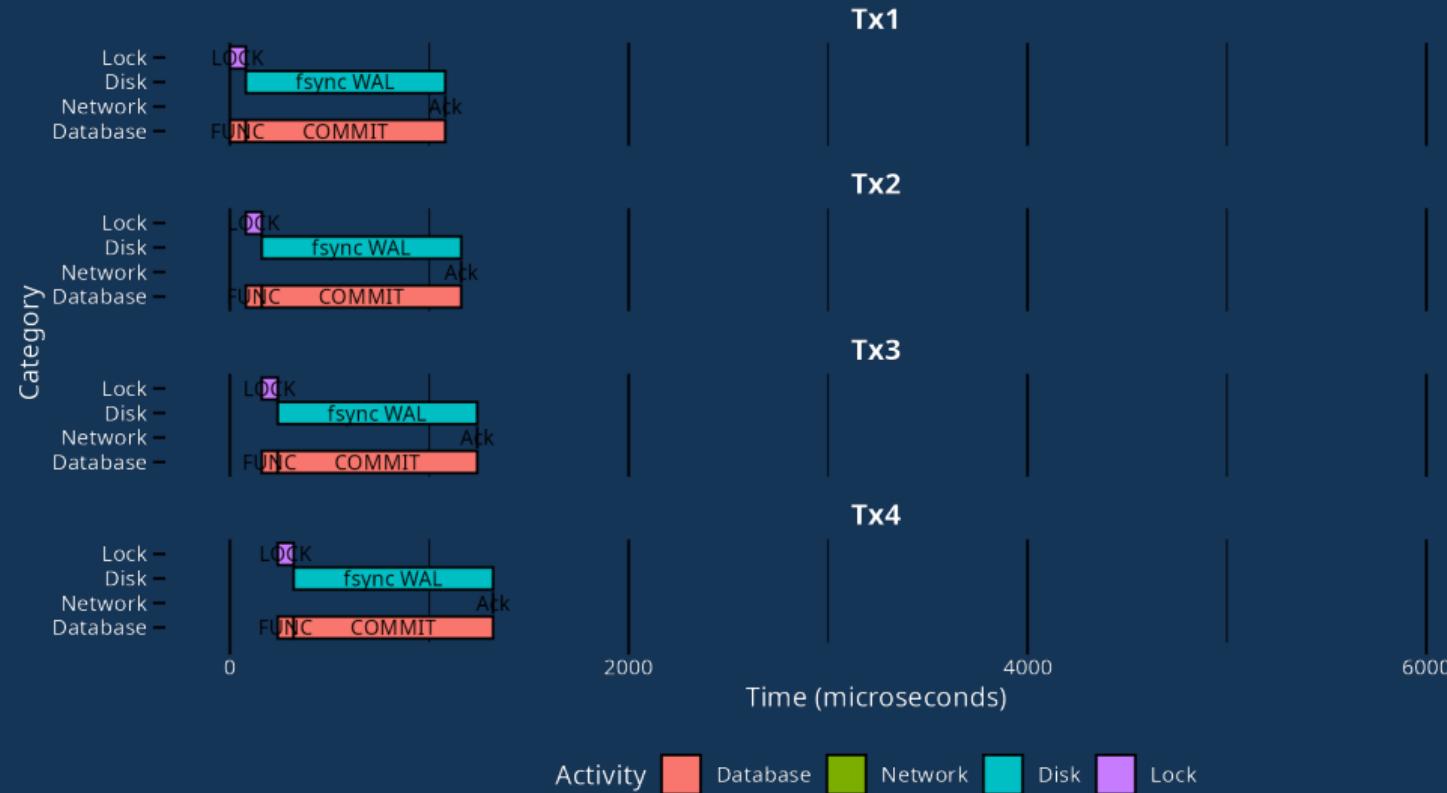
Eventual durability

- Write tx commit releases locks immediately.
 - Write becomes visible to new transactions.
- Write tx commit response waits for durability.
- Read tx keep track of newest tx seen.
- Read tx wait for durability of all seen tx on commit.



Eventual durability timeline

Transaction timeline with eventual durability



Eventual durability benefit

Almost synchronous_commit=off throughput on contended workloads with transactional guarantees.

No client side batching needed.



Prototype!

```
eventual durability$ pgbench -b plpgsql-tpcb-like -c96 -j96 -P1 -T300
pgbench (18.1)
starting vacuum...end.
progress: 1.0 s, 30405.4 tps, lat 3.097 ms stddev 1.305, 0 failed
progress: 2.0 s, 31759.0 tps, lat 3.026 ms stddev 0.981, 0 failed
progress: 3.0 s, 31589.9 tps, lat 3.029 ms stddev 0.968, 0 failed
progress: 4.0 s, 30789.1 tps, lat 3.130 ms stddev 1.151, 0 failed
progress: 5.0 s, 31343.9 tps, lat 3.062 ms stddev 1.034, 0 failed
progress: 6.0 s, 32102.1 tps, lat 2.992 ms stddev 0.942, 0 failed
progress: 7.0 s, 31409.0 tps, lat 3.050 ms stddev 1.152, 0 failed
progress: 8.0 s, 32696.1 tps, lat 2.940 ms stddev 0.930, 0 failed
```



Recap



What we learned today

- Avoid contended updates if you can.



What we learned today

- Avoid contended updates if you can.
- Figure out if READ COMMITTED is good enough.



What we learned today

- Avoid contended updates if you can.
- Figure out if READ COMMITTED is good enough.
- Minimize transaction length.



What we learned today

- Avoid contended updates if you can.
- Figure out if READ COMMITTED is good enough.
- Minimize transaction length.
- Get rid of network roundtrips using server side functions.



What we learned today

- Avoid contended updates if you can.
- Figure out if READ COMMITTED is good enough.
- Minimize transaction length.
- Get rid of network roundtrips using server side functions.
- Amortize durability delay by batching.



Thank you



Questions?

